




# DOCA 1.1 & BLUEFIELD OS V3.7 RELEASE

June 2021



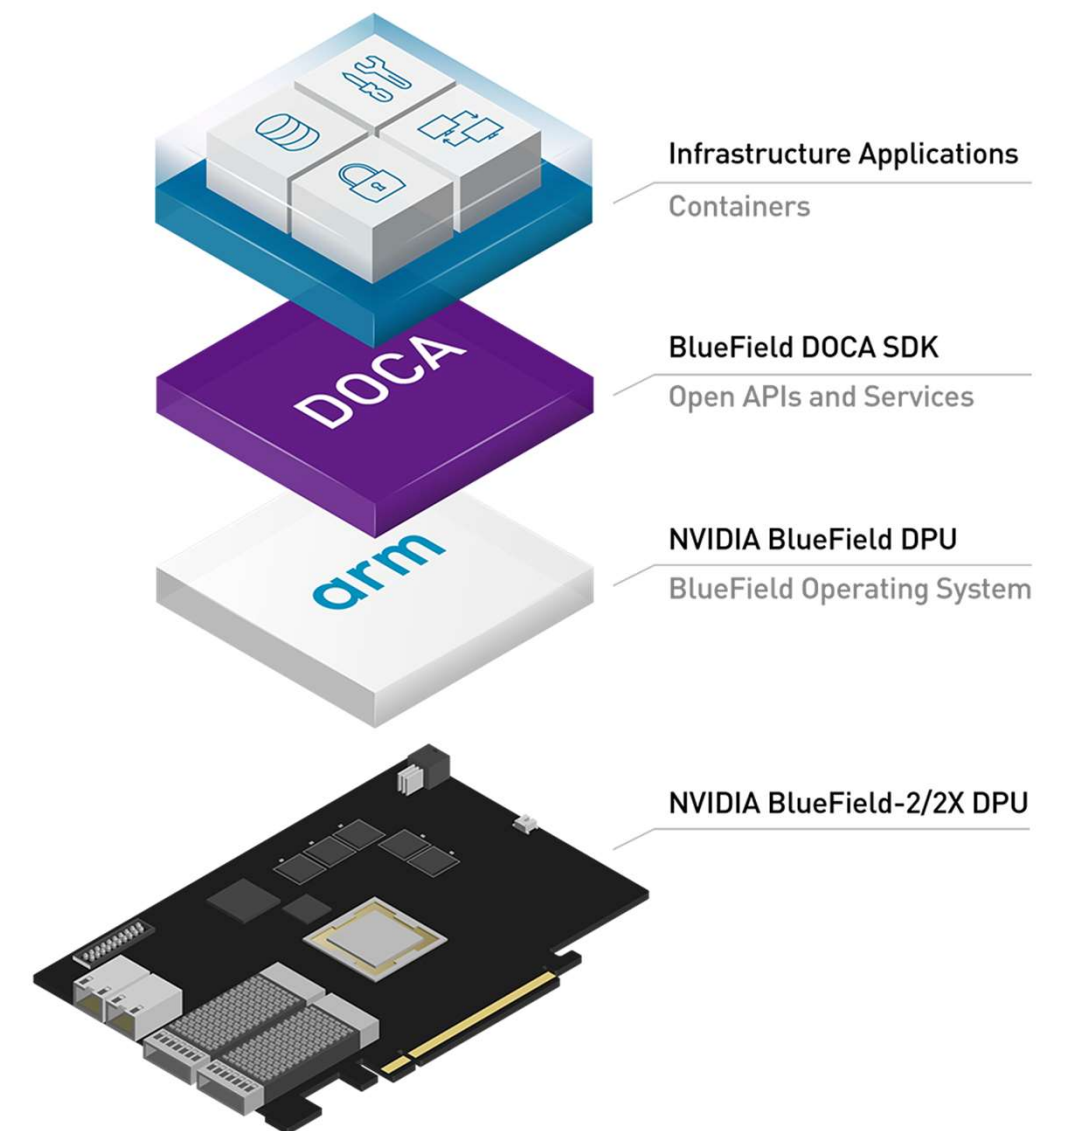
# DOCA 1.1 BLUEFIELD OS V3.7 GA

## Enabling the Transition to DPU based Data Center architecture

|              |  |
|--------------|--|
| Products     | DOCA v1.1 BlueField OS v3.7 GA<br>DOCA SDK v0.2  |
| Platforms    | BlueField-2 DPU 25G & 100G GA<br>BlueField-2 DPU Controller 100G ES  |
| Key Features | DOCA FLOW Lib - Gateway<br>DNS & URL Filter example applications<br>Acceleration of Stateful Flow Table (SFT)<br>DOCA DPI Lib with SFT & RegEx acceleration<br>DOCA Packages for x86<br>BlueField SNAP virtio-blk Resilience<br>BlueField SNAP Direct for virtio-blk |

Availability [Software download](#)

Actions Encourage customers to sign up for [DOCA early access](#)





# Features Highlight

## DOCA

Flow Library - Gateway

DPI Library

Reference Applications

DOCA Developer Experience

---

## Security

Secured NIC

TLS Inline Offload

---

## Storage

SNAP virtio-blk Resilience

SNAP Accelerated flow

SNAP Process Upgrade / crash recovery

SPDK NVMe-oF RDMA DC

---

## Supported OS & Cards

BlueField-2 3.7 Portfolio

Supported OS

---

## Enhancements





# DOCA SDK HIGHLIGHTED FEATURES

# DOCA FLOW LIBRARY - GATEWAY

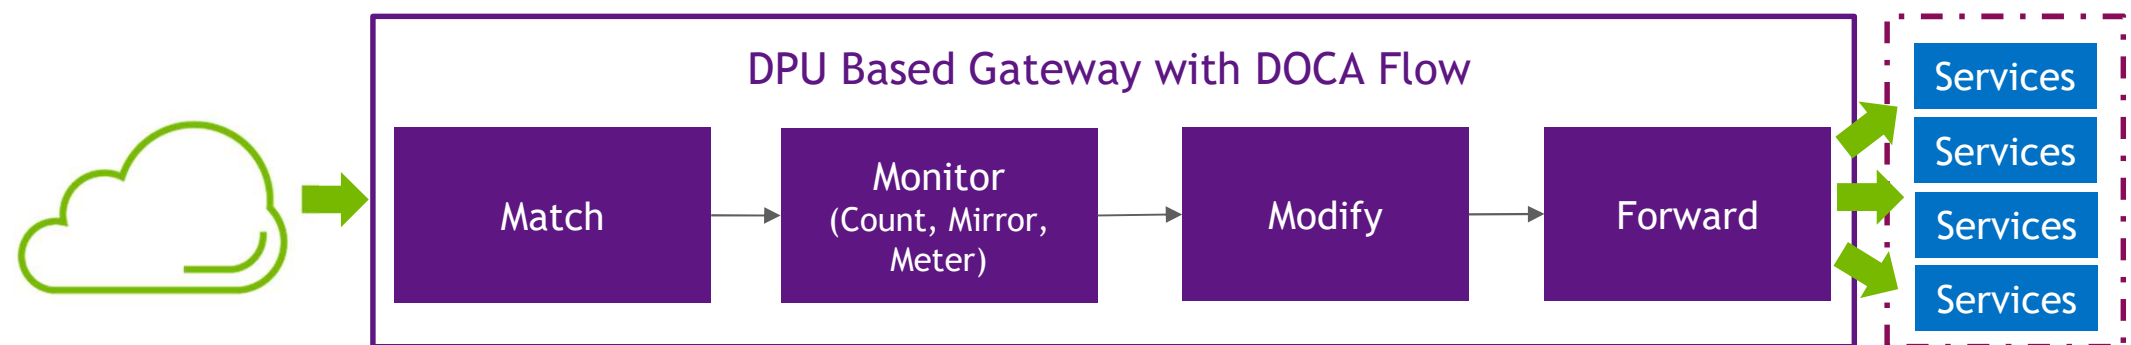
Creating Optimized and Accelerated Gateways

## Developer's Benefits

- Simplicity and scale through dedicated use case APIs and logic
- Tuned for best performance
- VNF and Appliance
- High insertion rate
- Supporting remote API calls and multi threading

## Use Cases


- Data network gateway
- Service Load balancer
- Carrier Grade NAT



# DOCA FLOW GATEWAY - DNS FILTERING

## Example Application

- Reduce CPU overhead by offloading DNS requests handling to the DPU Arm control plane
- DNS filter [example application](#) is leveraging DOCA Flow Gateway APIs to steer DNS requests and non-DNS traffic
- Further processing can be done on the DPU Arm such as whitelisting, logging, filtering, etc
- Non-DNS traffic is forwarded to the host or forwarded by the hardware using hair-pinning
- DNS filter example application is designed to run as a "bump-on-the-wire"



# DOCA DPI LIBRARY

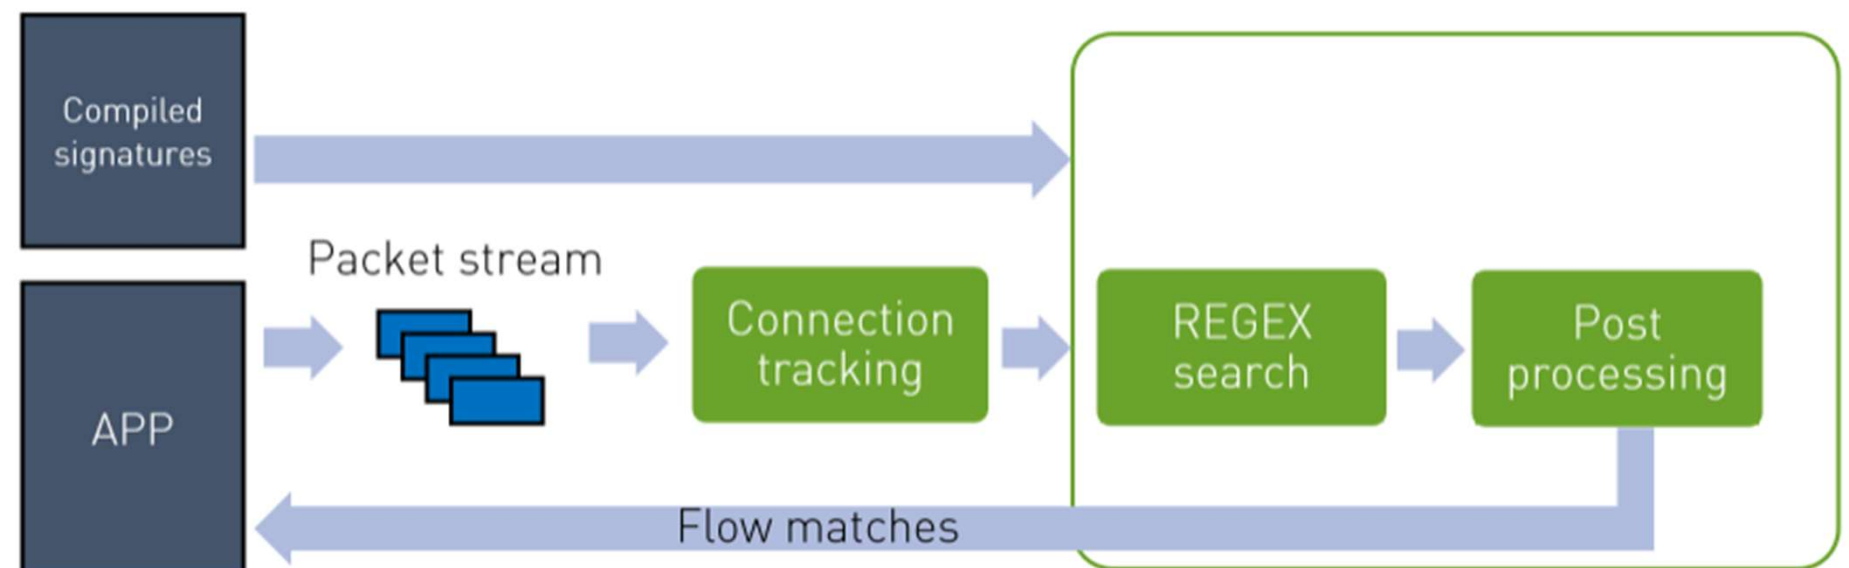
Enable accelerated & distributed Deep Packet Inspection

## Developer's Benefits

- Simplicity and scale through dedicated use case APIs and logic
- Tuned for best performance using SFT (Connection Tracking) and RegEx (Regular Expression) acceleration in hardware
- VNF and Appliance ready
- 100Ks of signatures are supported
- Total inspection bandwidth of 50Gbps

## Use Cases

- Next Generation Firewall acceleration
- Intrusion detection
- Stateful L7 load balancers
- Content policy violations

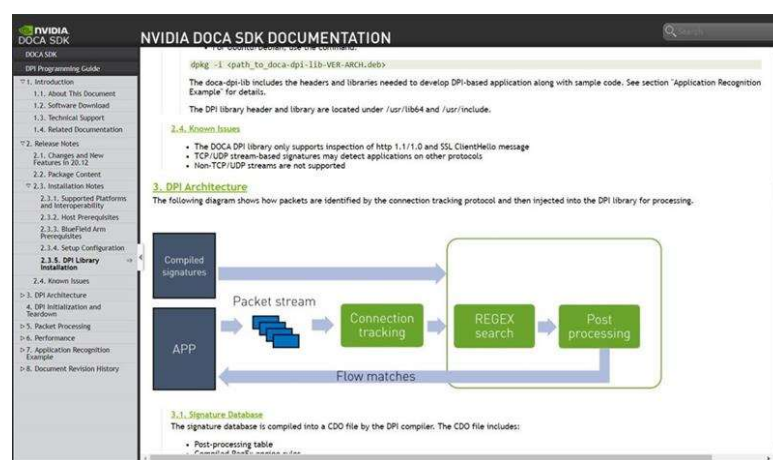


# REFERENCE APP—DPI FOR APPLICATION RECOGNITION

## Step-by-Step Developer Journey


1

Review the [DOCA Dev Zone](#) "Programming Guide" for Deep-Packet-Inspection (DPI) Library



2

Review the DOCA "Reference Applications" section for Application Recognition (AR)



Approach NVIDIA support for more information on precompiled signature databases

3

Implement changes on top of the reference application or use the DPI APIs from your application

DPI compiler is needed if you're using your own signatures





# REFERENCE APP—DPI FOR APPLICATION RECOGNITION

## Step-by-Step Developer Journey

4

Compile your application on the emulated Arm developer container or natively on the Arm

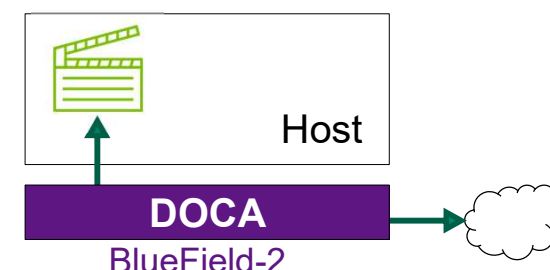


Load your application to the BlueField OS that has DOCA Runtime installed

5

Browse "Video App" from the host and make sure there's connectivity


Review the Application Recognition terminal output with App name



Use the NetFlow exporter to export the traffic statistics

6

Block "Video App" on the host by issuing block command per "Video App" ID on the DPU



Verify that Video App is now blocked

# DOCA DEVELOPER EXPERIENCE

## SDK Manager

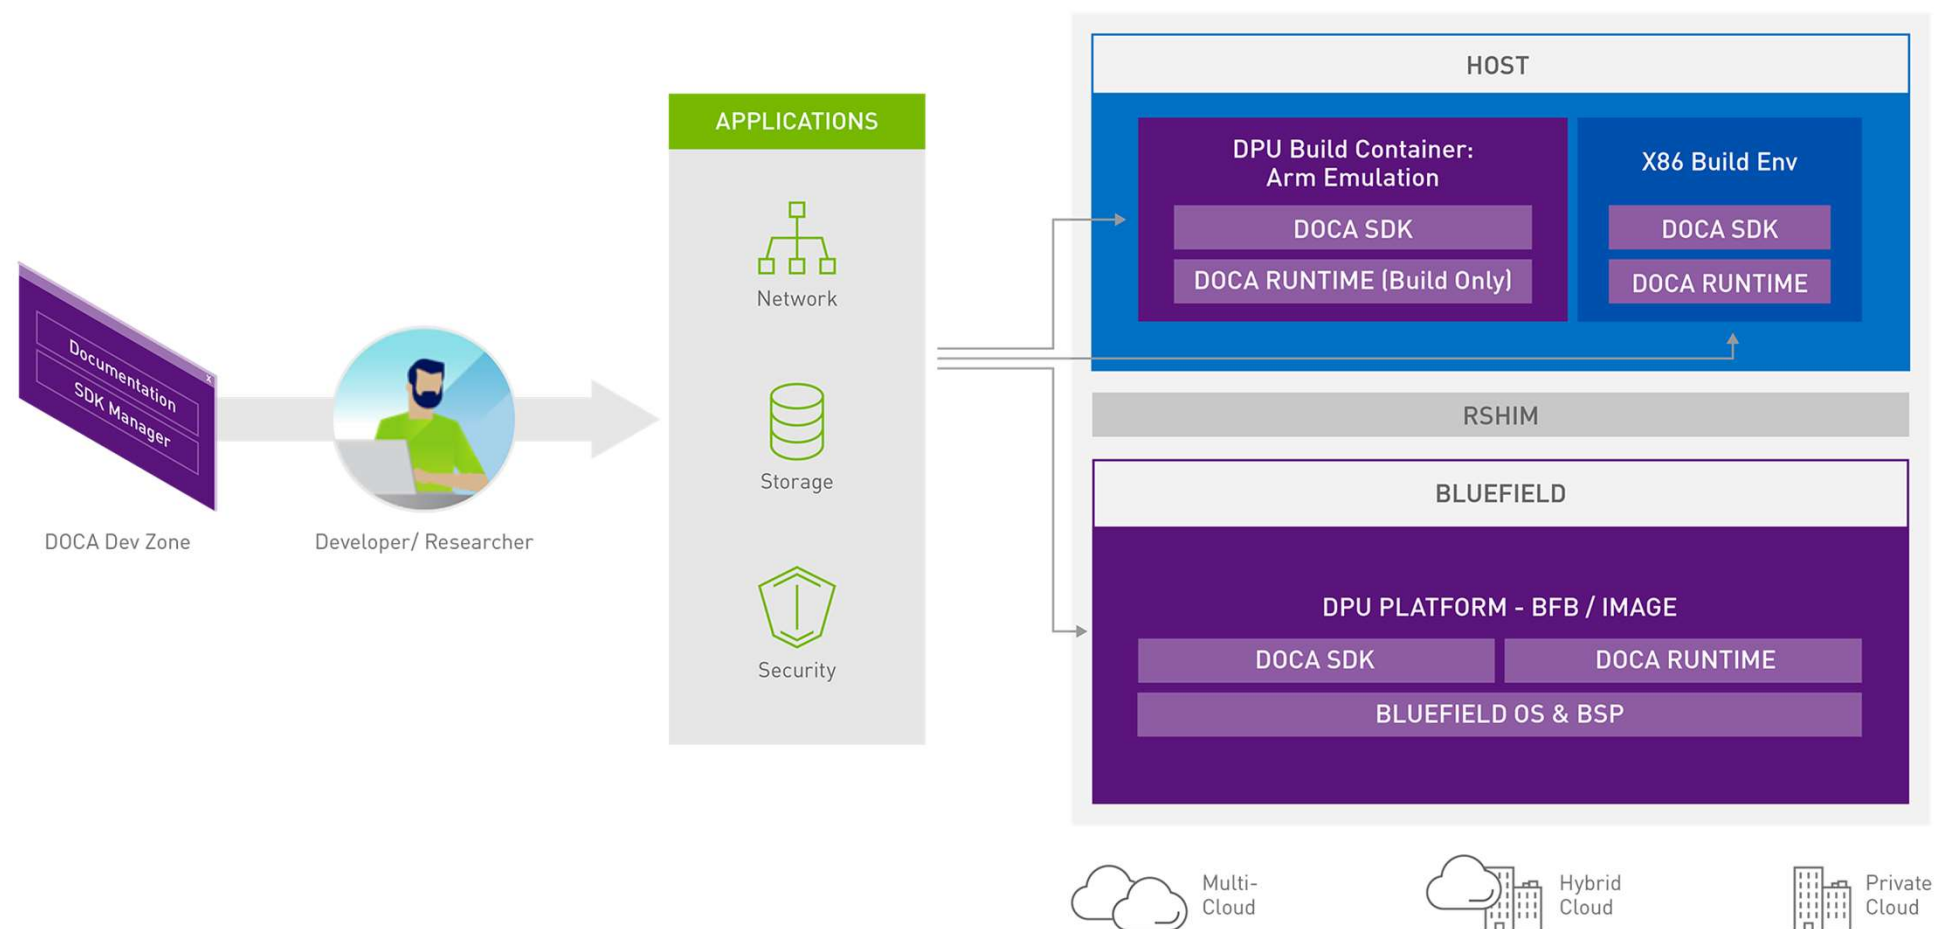
Easily install the DOCA packages  
Flash or Update DPU Software

## Developer Container

Develop DOCA anywhere  
Build services & images on host

## DOCA Package for x86

Develop host-based applications  
Leverage the DPU acceleration



A network graph visualization on a dark background. It consists of numerous small circular nodes connected by thin, light gray lines representing edges. The nodes are distributed across the frame, with a higher density in the upper half. Some nodes are highlighted in a bright yellow-green color, while others are a lighter, semi-transparent gray. The connections between nodes form a complex web, with some nodes acting as hubs with many incoming and outgoing edges.

# SECURITY HIGHLIGHTED FEATURES

# NIC FIRMWARE ROLLBACK PROTECTION

BlueField-2 includes 2 subsystems - NIC (ConnectX-6DX) & ARM

NIC subsystem now includes protection to enforce inability to downgrade FW version

Value - added protection (based on security version numbers in HW eFuses) from malicious actor rolling back to a FW with known vulnerabilities

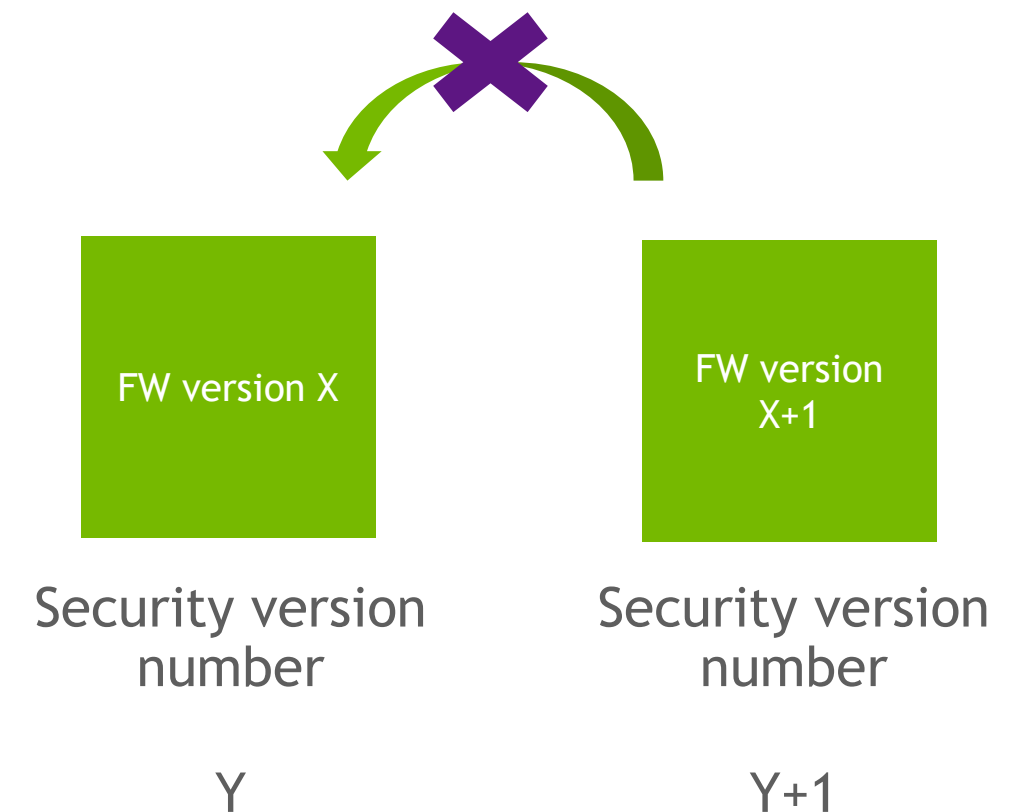
ARM subsystem already supports similar ability

Can be enabled on cards in field (with 32 security level increments) by updating to latest FW

"Fresh" New cards can have their "minimal version" fuses burned in factory

New field in flint query to provide visibility

Support for NIC side in BlueField-2 onwards



```
Secure boot: Enabled
EFUSE Security Ver: 0
Image Security ver: 0
Security Ver Program: Manually ; Disabled
```

# TLS INLINE OFFLOAD LAG SUPPORT

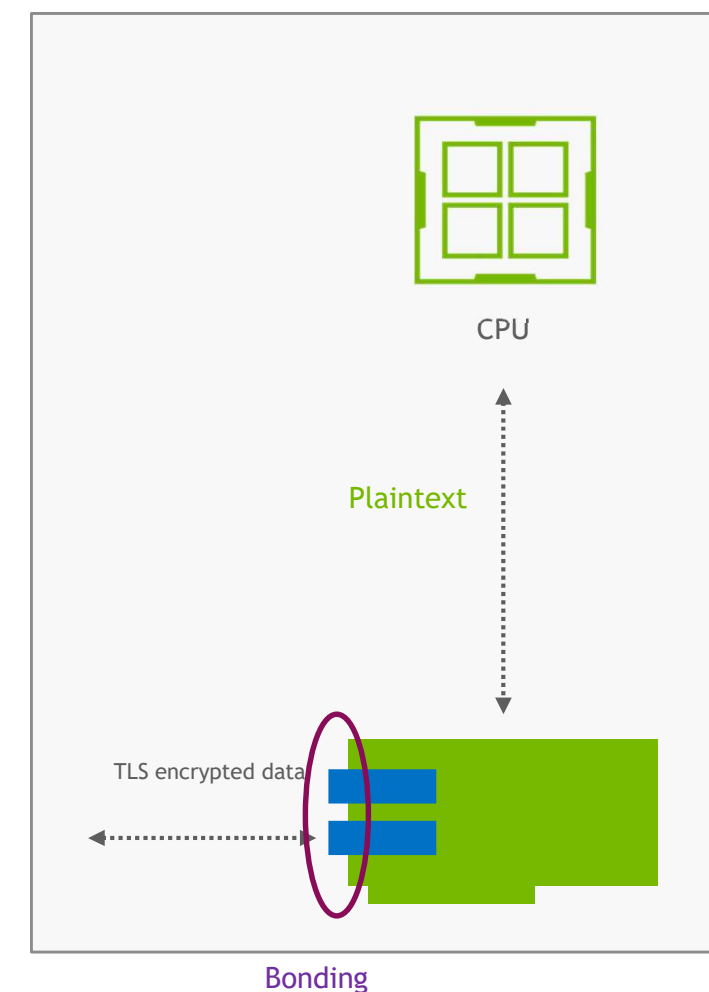
Support link aggregation groups (LAG) with kTLS crypto HW offload

Aggregating multiple network interfaces into a single logical bonded interface

Value -

Allows maximizing aggregated bandwidth in relevant system setups with multiple network interfaces

Enables redundancy in case one of the links should fail





A network graph visualization on a dark gray background. The graph consists of numerous small, semi-transparent nodes. Some nodes are white, while others are a bright yellow-green. These nodes are interconnected by a dense web of thin, light gray lines, representing connections or edges. The overall structure is complex and non-linear, with some nodes appearing to be central hubs with many connections, while others are more isolated. The text 'STORAGE HIGHLIGHTED FEATURES' is overlaid in the bottom right corner in a white, sans-serif font.

STORAGE  
HIGHLIGHTED FEATURES

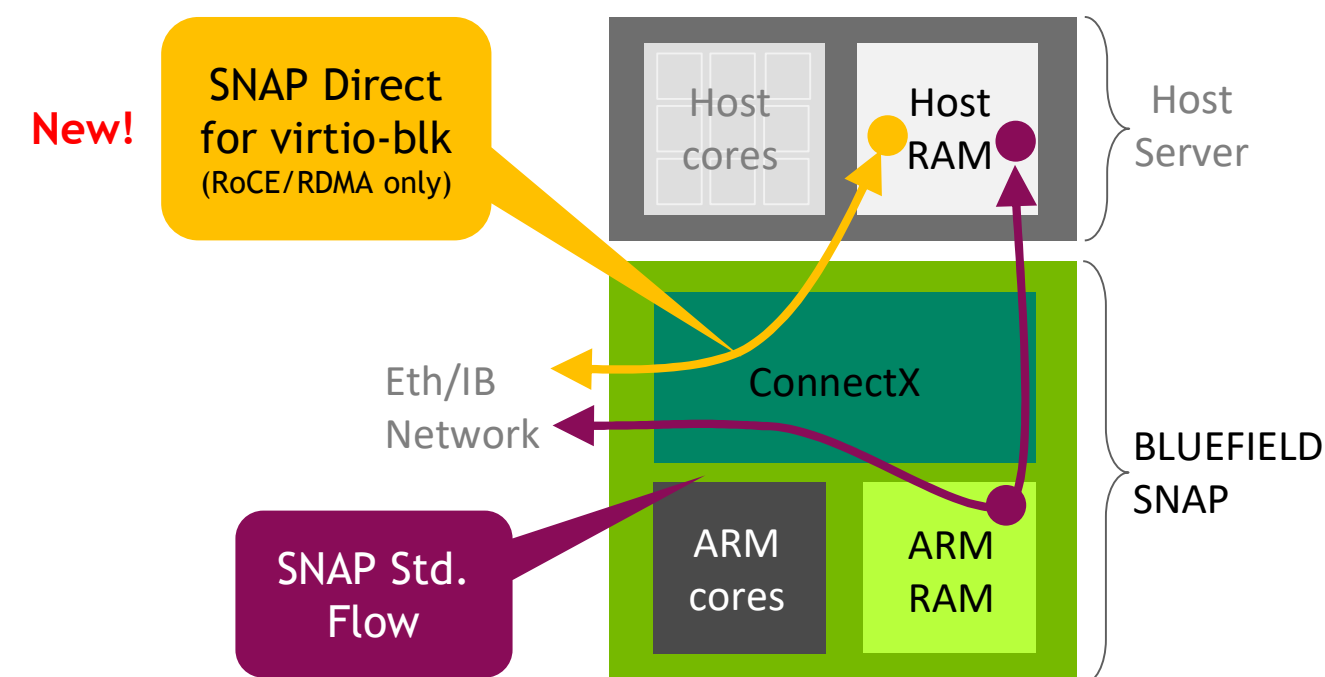
# virtio-blk SNAP-DIRECT

Lowers latency; frees Arm memory BW;  
increased efficiency

- SNAP-Direct introduced in previous release for NVMe only
- New in this release - SNAP-Direct for virtio-blk

## Capabilities

- Direct Host RAM buffer to Network tx/rx (zero-copy)
- Control remains on Arm cores
- Currently supported for RoCE/RDMA only - NVMe-oF is main use case



# SNAP ACCELERATED FLOW

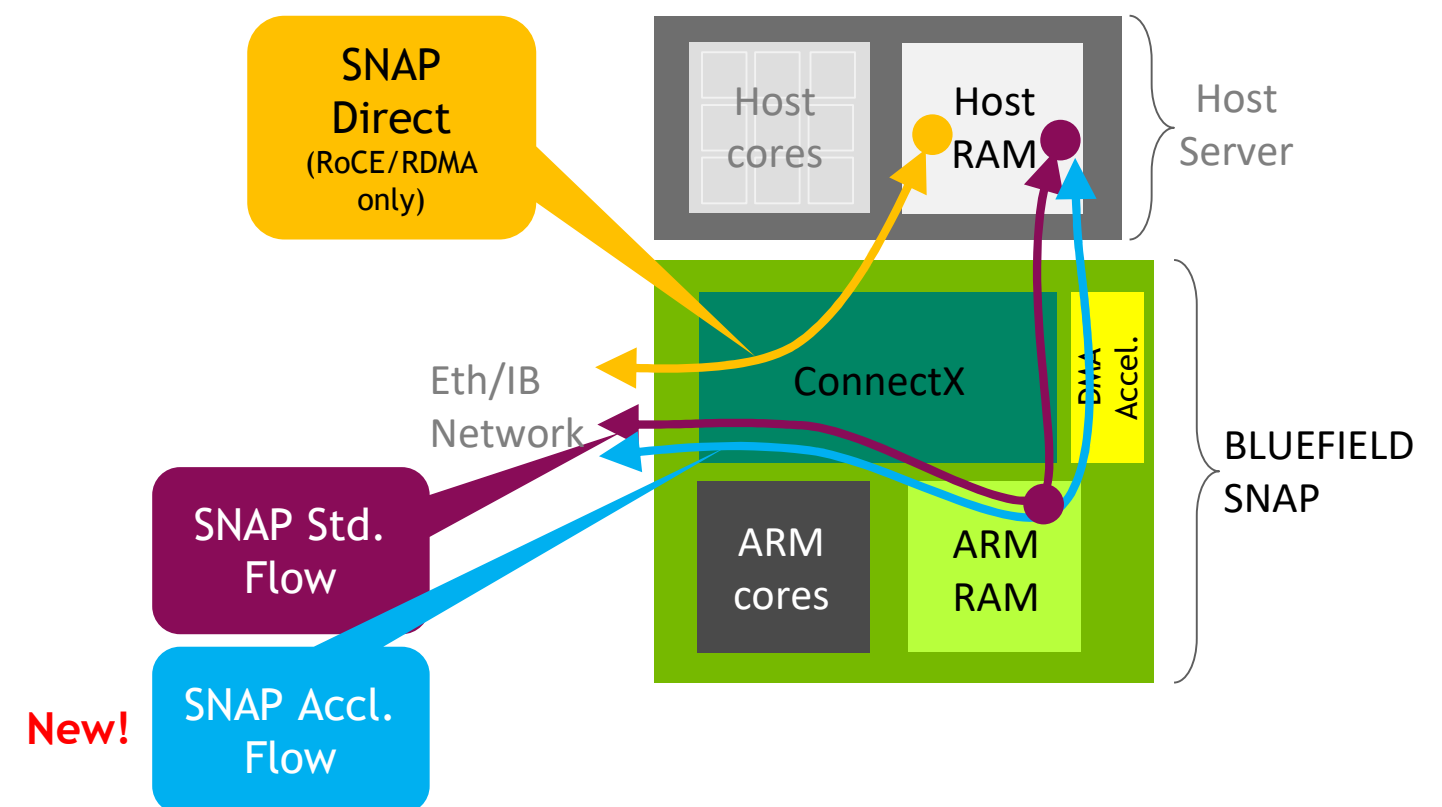
## Accelerated ARM-DATA-PATH

Improving 'SNAP Std. Flow' using new-BlueField-2-DMA engine

Data-path Host-ARM is accelerated, hopping PCIe only once

DMA-engine is a 400Gbps device (200Gbps full-duplex)

Performance with data copy to ARM RAM is expected to be ~4MIOPS @4K blocks



# SNAP PROCESS UPGRADE / CRASH RECOVERY



## Software Upgrade

- Allow SNAP/SPDK process upgrade during traffic
- Run next-ver and recover traffic from same point
- Quick to recovery
- External Orchestration dependent
- virtio-blk only



## Crash Recovery

- In the event of SNAP/SPDK process-crash
- Restart to recover traffic from same point
- Quick to recovery
- External Orchestration dependent
- virtio-blk only

# SPDK NVME-OF ADDS RDMA DC SUPPORT



Today, most widely RDMA Transports used are -

- UD (Unreliable Datagram)

- ✓ **Unreliable**

- ✓ **Scalable: One QP services multiple destinations**

- RC (Reliable Connection)

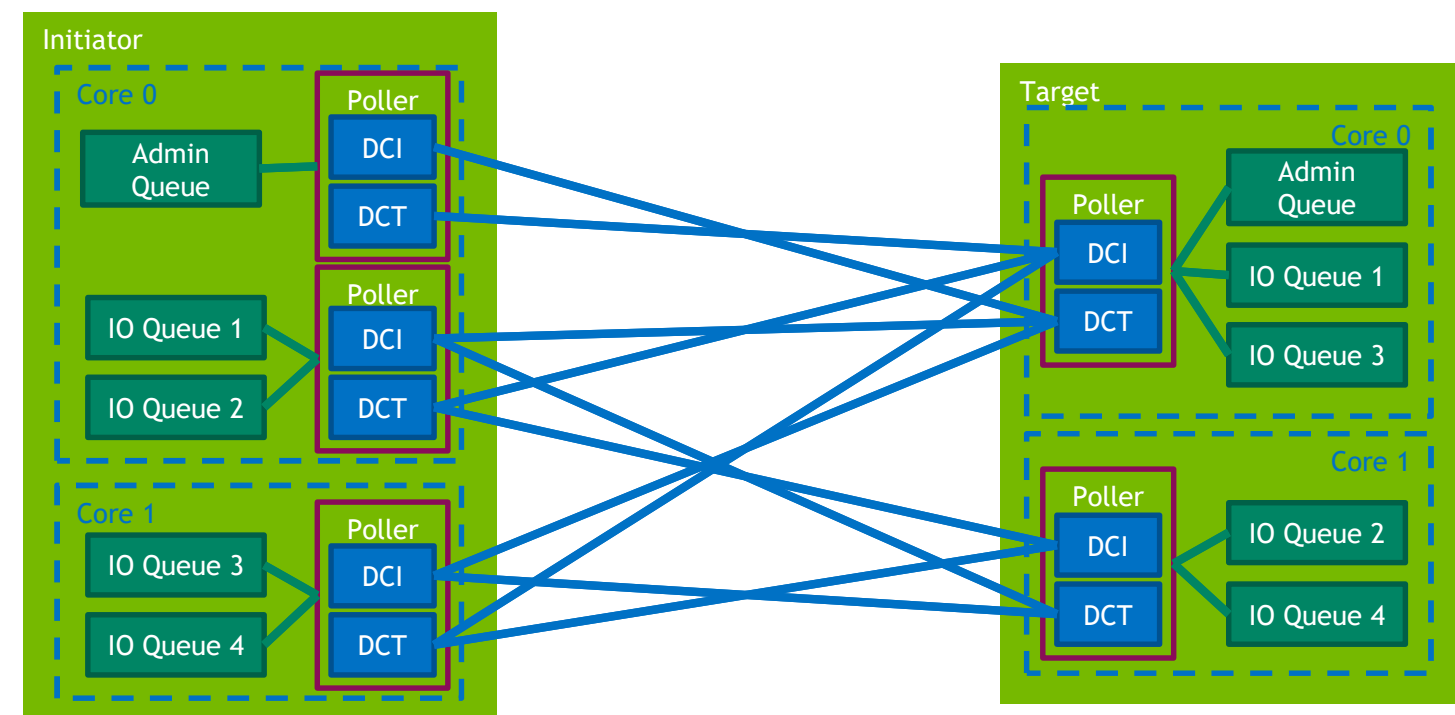
- ✓ **Reliable**

- ✓ **Scalability limitations: One connection per destination**

RDMA DC Transport provides best of both worlds -

- ✓ **Reliable**

- ✓ **Scalable: One QP services multiple destinations**



- SPDK adds RDMA DC support
- Initiator and Target (DCI/DCT)
- Integrated into NVMe-oF for both RoCE (Eth) and RDMA (IB)
- Beta Available now (private ver); Upstreaming is on-going






SUPPORTED OS & CARDS

# BLUEFIELD-2 3.7 PORTFOLIO

FOR ALIBABA ONLY



BLUEFIELD-2 DPU  
25G End Point

GA



BLUEFIELD-2 DPU  
100G End Point

GA



BLUEFIELD-2 DPU  
CONTROLLER  
100G Root Complex

Sample

FOR ALIBABA ONLY

|         |                             |   |                               |
|---------|-----------------------------|---|-------------------------------|
| Flavors | 2x25G<br>P-series (2.5 GHz) | 2x100G<br>E-series (2.0 GHz)<br>P-series (2.75 GHz) | 2x100G<br>P-series (2.75 GHz) |
|---------|-----------------------------|---|-------------------------------|

# SUPPORTED OS

FOR ALIBABA ONLY

## DPU ARM OS

Ubuntu Server 20.04 Kernel 5.4 [default]

CentOS 7.6 (drivers only for LK 5.4) \*  
CentOS 8.2 (drivers only native LK and LK 5.4) \*

Debian 10 (drivers only) \*

## HOST OS

Ubuntu 16.04 \* / 18.04 / 20.04

RHEL/CentOS 7.4 \* / 7.5 \* / 7.6 / 8.0 / 8.2

Debian 9.11 \*

Windows Server 2016 \*, 2019 \*

\* Not supported in DOCA, please approach the support team



Enhancements

# DOCA 1.1 BLUEFIELD OS V3.7 ENHANCEMENTS

## DOCA SDK v0.2

x86 DOCA Packages - Runtime and SDK packages for x86

## Platform

NVSM Integration

Reduce Boot time

DOCA Telemetry Service Early Access

## Storage

Signature Offload API



